

RecruitMe — Project Scope, Current State & Running Costs

*Prepared 2026-06-04; last updated 2026-06-22 — adds **Pulse** (SEEK CV-update alerts), **per-user activity** + **AI-cost/failure tracking**, and **inline-editable candidate contact details & links**, on top of the 2026-06-15 production activation of CRM, tags + reminders, white-label theming, identity merge-review, and the insight-ranked talent flywheel. Figures from live production data where measured; estimates flagged.*

1. What it is

RecruitMe is a recruiter-facing SaaS that sources, scores, and manages candidates for open roles. A recruiter creates a **job** (pastes a JD → AI parses it into structured must-haves / nice-to-haves / seniority / location / salary), then **finds candidates** from three sources — an internal **talent pool** (~15k profiles), **LinkedIn**, and **SEEK Talent Search** — and the system **AI-scores** each candidate against the role, producing a ranked shortlist with a written rationale per candidate.

Multi-tenant (org-scoped); currently one organisation (placeMe IT) with a 5-person recruiter team. Deployed on Railway; live discovery runs through a self-hosted mini-PC scraper.

2. Current state (what's built & working)

Scale (live, 2026-06-22): 16,284 candidates · 983 AI-scored · 827 with AI-extracted insight-facts · ~13.5k CVs stored (encrypted at rest) · 14 jobs · 5 active Pulse watches (76 profile-update hits surfaced) · 5 users in 1 org.

Tech stack

- **App:** Next.js 15 / React 19 / TypeScript, Prisma ORM, PostgreSQL 18 — hosted on **Railway** (auto-deploy from GitHub `main`).
- **AI:** Anthropic Claude — **Haiku 4.5** for parsing & snippet scoring (cheap), **Sonnet 4.6** for full-profile scoring (quality). Cost-tracked per call with a daily spend cap.
- **Scraper:** self-hosted mini-PC ("the box") running a persistent logged-in browser (Patchright) for LinkedIn + SEEK; reached over Tailscale. Single-worker by design (ban-avoidance).
- **Storage:** CVs in an S3-compatible bucket (`t3.storageapi.dev`), **AES-256-GCM encrypted at rest**, original sender format preserved.

Features shipped & working

- **JD analysis** — Claude parses a pasted JD into structured requirements + search expansion + title variants.
 - **Multi-source candidate search** — internal pool (Postgres full-text + AI ranking), live LinkedIn scrape, live SEEK Talent Search scrape. Library-first (cheap pool results instant; live sources on demand).
 - **AI scoring** — per-candidate match score (0–100) with must-have coverage, category breakdown, and a recruiter summary; provisional (snippet) vs full-profile scoring; per-org configurable weights; cached so unchanged candidates aren't re-billed.
 - **SEEK Talent Search** — ingests result cards as candidates (credit-safe: harvests free card data, no per-profile credit spend), now **sorted by "Date updated"** so the freshest movers lead the results.
- Validated 2026-06-04; recency sort 2026-06-17.*

- **LinkedIn discovery** — people-search → profiles scraped → ingested + scored. (Fixed 2026-06-15: a stale hardcoded NZ geo-filter parameter that LinkedIn had silently stopped honouring was making **every** LinkedIn search return "No results found" — verified on the box, then removed. Results are NZ-biased via the logged-in account; the app re-narrows by the requested location.)
- **CRM — clients / submissions / placements (live)** — client records, "submit candidate to client" from a job, placement tracking with fee/guarantee/invoice + auto-reminders. Finishes the previously half-built CRM; all create paths org-FK-validated.
- **Candidate tags + reminders (live)** — colour-coded tags (create / assign / manage at /settings/tags) and a reminders bell with due/overdue tracking (follow-up · guarantee-check · client-feedback · custom).
- **White-label theming (live)** — per-org brand colour (inert until an org sets one; computes a distinct hover shade + luminance-darkens too-light colours so button text stays readable).
- **Identity merge-review (live)** — recruiter-facing panel on the candidate page to confirm or split duplicate-person merges (same email / phone / name surfaced across sources), over the org-scoped, tombstone-safe merge/unmerge routes.
- **Talent pool / flywheel (live, insight-ranked)** — every search enriches the pool; future searches serve those candidates instantly. **Insight-ranked search is now on:** 812 candidate profiles carry AI-extracted structured facts (primary stack / titles held / domains) that nudge library-search ranking — a bounded re-rank (own-org-scoped, a no-op for any candidate without an insight, never reorders across the FTS top results by more than a few places).
- **Candidate management** — library, per-job pipelines (shortlist/contacted/etc.), CV upload + text extraction, identity dedup across sources (LinkedIn / SEEK / JobAdder), screening data, recruiter notes, CSV export.
- **CV handling** — extraction from PDF/Word/RTF, encryption at rest, original-format preservation.
- **Ops** — on-box health dashboard (CPU/RAM/disk/temp, scraper status), cost attribution + daily spend cap, PII redaction on the seller heartbeat, CI real-DB test gate, error reporting.
- **Security & data-safety floor (shipped 2026-06-15)** — CV encryption now fails closed (refuses to mint a new key when encrypted CVs exist, so a missing key can't silently orphan them); cross-org candidate-tag read/write leak closed; four AI-cost-attribution gaps closed (spend now always billed to the right org); Sentry PII scrubber so candidate names/emails/profile text don't leave the app; Llama-offload default flipped off (the box can't score, so a credit-out now fails cleanly instead of hanging).
- **Live JobAdder feed** — the operator's JobAdder candidate data (~13k profiles) is available live and on demand, drawn from the operator's own JobAdder account, and kept current in the platform.
- **Pulse — SEEK profile-update alerts (live)** — a recruiter defines a "watch" (boolean query + location + run interval); the box re-runs it on SEEK on a schedule and surfaces candidates who've **updated their SEEK profile** into a live, terminal-style feed and the reminders bell, with the profile linked. Credit-safe (card harvest only), recency-sorted (newest movers first), and each watch/hit shows who set it up. A box timer drives the schedule; the watch always notifies from when it's created.
- **Per-user activity + AI-cost/failure tracking (live)** — the admin dashboard attributes every action (search / score / parse / capture) to the user who did it and which job, with a per-user breakdown — plus an **AI-failure view** so the operator can see who hit an out-of-credit / rate-limit wall, and when. Built for the now-multi-user team.
- **Inline-editable candidate fields (live)** — LinkedIn / JobAdder / SEEK links and phone / email / location are now editable directly on the candidate page (add, correct, or clear), each with a matching brand icon. Previously these were display-only.

Known constraints / debt

- **Two parallel search systems** — now largely converged onto one durable, resumable multi-source run (job-scoped, live-streaming). The legacy job-search path is mostly retired; final dead-code cleanup is the remaining item — see roadmap §3.
- SEEK candidates are **snippet-level** (name/role/company/location); full-profile enrichment is a separate, credit-gated action (scraper exists, not auto-fired). (*Fixed 2026-06-15: "Score all" now scores these thin LinkedIn/SEEK finds as a low-confidence estimate instead of silently skipping them.*)
- ~16.3k candidates, **983 scored** — large unscored backlog. Scoring is **manual + on-demand** by design (recruiter clicks *Score all* or per-card) and cost-capped; finds are not auto-scored on import/discovery. (*Score-all now auto-resumes if a long run's connection drops, picking up where it left off rather than restarting.*)
- **No AI fallback** — production is 100% Claude-dependent (no working local model on Railway or the box). If the Anthropic credit balance runs out or Claude is unavailable, all AI — scoring **and** insight extraction — hard-fails with no graceful degradation. Operationally: keep the Anthropic balance topped up.
- Single scraper worker (no redundancy); the box must stay online for live discovery.
- Scraper still authenticates via a shared secret (org-binding token cutover incomplete).
- **SEEK session needs periodic manual re-login** — SEEK logs in through Auth0 + a Cloudflare Turnstile human-check and actively expires automated sessions, so the box's SEEK login has to be refreshed by a person every day or two (it can't be automated past the human check). The app softens this — a self-healing auth circuit and a rolling session re-save stretch the interval and auto-recover the instant a valid session returns — but can't fully remove it. Live SEEK search + Pulse pause until it's refreshed (one click in the on-box control dashboard). A fix to persist the Auth0 token for longer-lived sessions is in progress.

3. Future implementations (roadmap)

The authoritative, sequenced plan to take the app to feature-complete + self-running lives in `docs/completion-plan.md` (7 phases, built from a verified 6-dimension audit). The items below are the strategic summary.

A. Search architecture consolidation (*planned, documented in `docs/search-architecture-plan.md`*) — collapse the two parallel search systems into one **federated dispatcher + pluggable source adapters** (pool/LinkedIn/SEEK/+future), one durable result model, and a **live-streaming results UI** (results appear as they're scraped). Makes adding a new source trivial and removes the current fragility. ~6 phases, each independently shippable.

B. SEEK full-profile enrichment — one-click, credit-aware deep-fetch of a chosen SEEK candidate's full profile (the scraper exists; needs wiring + a credit guard).

C. Scoring at scale — score the unscored backlog; background/batch scoring; richer re-score-on-change.

D. Talent flywheel — later phases — the **read-path (insight-ranked search) is now live + backfilled (812 profiles)**. Remaining: background discovery that continuously enriches the pool (cost-gated); profile freshness / re-fetch of stale profiles; raising the insight coverage as the pool grows.

E. Candidate UX polish — **shipped** (email/screening/file-count surfacing, import-batch grouping, source badges, plus **inline-editable LinkedIn/JobAdder/SEEK links + phone/email/location** and matching brand icons). Remaining: candidate-card density declutter.

F. Hardening — complete the scraper org-binding token cutover; **persist the SEEK Auth0 token so the session auto-renews** (removes the periodic manual re-login — see constraints §2); multi-worker /

redundancy if scrape volume grows; observability.

G. SEEK credit visibility (*scoped, parked*) — surface the SEEK credit balance + burn-rate in-app (the balance is readable from SEEK's usage page; searches are credit-free, so the meter tracks profile-reveal spend). Awaiting go-ahead on depth.

Priority order is the operator's call; A unblocks the most future work.

4. Running costs per month

Item	Cost / month	Basis
Claude API (Anthropic)	~\$15–40	Mostly cheap Haiku (parsing, snippet scoring, insight extraction). With Sonnet scoring full profiles (\$0.05/score) cost rises with scoring volume; hard-capped at \$5/day (~\$150/mo ceiling) by the spend guard. Realistic steady state ~\$20–40 . One-off: the talent flywheel insight backfill (812 profiles, Haiku) cost **\$4** total. ⚠️ No fallback — if this balance hits zero, all AI stops (see constraints §2).
Railway (app + Postgres + egress)	\$25	Hosting: always-on Next app + Postgres (5 GB volume) + bandwidth. Operator-confirmed figure.
CV storage (S3 / t3.storageapi.dev)	~\$1–3	13.5k encrypted CVs (8–12 GB est.) at S3-compatible rates + light egress.
Mini-PC scraper (electricity)	~\$3–6	Low-power i3 mini-PC 24/7 at NZ power rates. Internet = existing home line; Tailscale = free tier.
Domain	\$0	Railway subdomain (*.up.railway.app). Custom domain would add ~\$1–2/mo.
Platform infra subtotal	~\$45–75 / month	Everything RecruitMe itself costs to run (recurring).
Mini-PC hardware (one-off)	\$150 once	The mini-PC itself — already purchased; one-time capital cost, not recurring (≈ \$4–6/mo amortised over 2–3 yrs).
SEEK Talent Search	<i>operator's existing subscription</i>	External recruiting tool (the box logs into the operator's SEEK account). Significant but pre-existing; per-search/profile credits apply — the app is built to be credit-safe.
LinkedIn	<i>operator's existing account</i>	Scraper uses the operator's logged-in LinkedIn; no extra app cost.

Bottom line: the RecruitMe *platform* runs at roughly **\$45–75/month** in direct infrastructure + AI cost (plus a one-off **\$150** mini-PC, already purchased), on top of the operator's existing SEEK/LinkedIn/JobAdder subscriptions. The single largest and most variable line is the Claude API, which is bounded by the daily spend cap and scales with how much scoring is done.

5. Tech-stack summary (for a developer handoff)

- **Frontend/Backend:** Next.js 15 (App Router), React 19, TypeScript, Tailwind.
- **Data:** PostgreSQL 18 + Prisma; Postgres full-text search (FTS) for the talent pool.
- **AI:** Anthropic SDK (Claude Haiku 4.5 + Sonnet 4.6); structured scoring pipeline with caching + cost tracking + spend caps.
- **Scraping:** Node/TypeScript worker on a mini-PC, Patchright (patched Chromium), one persistent session per platform; polls the app's job queue (Postgres-backed `ScrapeJob` , priority + dedup).
- **Infra:** Railway (CI real-DB gate on deploy), Tailscale (box networking), S3-compatible blob storage (encrypted CVs).
- **Repo:** GitHub `baeley-canning/RecruitMe` , deploy = push to `main` .