

# RecruitMe — Project Scope, Current State & Running Costs

*Prepared 2026-06-04; last updated 2026-06-15 — **all previously flag-gated features are now activated in production** (CRM, candidate tags + reminders, white-label theming, identity merge-review, and the insight-ranked talent flywheel). Figures from live production data where measured; estimates flagged.*

---

## 1. What it is

---

RecruitMe is a recruiter-facing SaaS that sources, scores, and manages candidates for open roles. A recruiter creates a **job** (pastes a JD → AI parses it into structured must-haves / nice-to-haves / seniority / location / salary), then **finds candidates** from three sources — an internal **talent pool** (~15k profiles), **LinkedIn**, and **SEEK Talent Search** — and the system **AI-scores** each candidate against the role, producing a ranked shortlist with a written rationale per candidate.

Multi-tenant (org-scoped), single primary operator today. Deployed on Railway; live discovery runs through a self-hosted mini-PC scraper.

---

## 2. Current state (what's built & working)

---

**Scale (live, 2026-06-15):** 15,421 candidates · 731 AI-scored · 14,709 with captured profiles · 812 with AI-extracted insight-facts (across 1,513 dedup'd identities) · ~13.5k CVs stored (encrypted at rest) · 11 active jobs.

### Tech stack

- **App:** Next.js 15 / React 19 / TypeScript, Prisma ORM, PostgreSQL 18 — hosted on **Railway** (auto-deploy from GitHub `main`).
- **AI:** Anthropic Claude — **Haiku 4.5** for parsing & snippet scoring (cheap), **Sonnet 4.6** for full-profile scoring (quality). Cost-tracked per call with a daily spend cap.
- **Scraper:** self-hosted mini-PC ("the box") running a persistent logged-in browser (Patchright) for LinkedIn + SEEK; reached over Tailscale. Single-worker by design (ban-avoidance).
- **Storage:** CVs in an S3-compatible bucket ( `t3.storageapi.dev` ), **AES-256-GCM encrypted at rest**, original sender format preserved.

## Features shipped & working

- **JD analysis** — Claude parses a pasted JD into structured requirements + search expansion + title variants.
- **Multi-source candidate search** — internal pool (Postgres full-text + AI ranking), live LinkedIn scrape, live SEEK Talent Search scrape. Library-first (cheap pool results instant; live sources on demand).
- **AI scoring** — per-candidate match score (0–100) with must-have coverage, category breakdown, and a recruiter summary; provisional (snippet) vs full-profile scoring; per-org configurable weights; cached so unchanged candidates aren't re-billed.
- **SEEK Talent Search** — now ingests result cards as candidates (credit-safe: harvests free card data, no per-profile credit spend). *Validated 2026-06-04.*
- **LinkedIn discovery** — people-search → profiles scraped → ingested + scored. *(Fixed 2026-06-15: a stale hardcoded NZ geo-filter parameter that LinkedIn had silently stopped honouring was making every LinkedIn search return "No results found" — verified on the box, then removed. Results are NZ-biased via the logged-in account; the app re-narrows by the requested location.)*
- **CRM — clients / submissions / placements (live)** — client records, "submit candidate to client" from a job, placement tracking with fee/guarantee/invoice + auto-reminders. Finishes the previously half-built CRM; all create paths org-FK-validated.
- **Candidate tags + reminders (live)** — colour-coded tags (create / assign / manage at `/settings/tags` ) and a reminders bell with due/overdue tracking (follow-up · guarantee-check · client-feedback · custom).
- **White-label theming (live)** — per-org brand colour (inert until an org sets one; computes a distinct hover shade + luminance-darkens too-light colours so button text stays readable).
- **Identity merge-review (live)** — recruiter-facing panel on the candidate page to confirm or split duplicate-person merges (same email / phone / name surfaced across sources), over the org-scoped, tombstone-safe merge/unmerge routes.
- **Talent pool / flywheel (live, insight-ranked)** — every search enriches the pool; future searches serve those candidates instantly. **Insight-ranked search is now on:** 812 candidate profiles carry AI-extracted structured facts (primary stack / titles held / domains) that nudge library-search ranking — a bounded re-rank (own-org-scoped, a no-op for any candidate without an insight, never reorders across the FTS top results by more than a few places).
- **Candidate management** — library, per-job pipelines (shortlist/contacted/etc.), CV upload + text extraction, identity dedup across sources (LinkedIn / SEEK / JobAdder), screening data, recruiter notes, CSV export.
- **CV handling** — extraction from PDF/Word/RTF, encryption at rest, original-format preservation.
- **Ops** — on-box health dashboard (CPU/RAM/disk/temp, scraper status), cost attribution + daily spend cap, PII redaction on the seller heartbeat, CI real-DB test gate, error reporting.
- **Security & data-safety floor (shipped 2026-06-15)** — CV encryption now fails closed (refuses to mint a new key when encrypted CVs exist, so a missing key can't silently orphan them); cross-org candidate-tag read/write leak closed; four AI-cost-attribution gaps closed (spend now always billed to the right org); Sentry PII scrubber so candidate names/emails/profile text don't leave the app; Llama-offload default flipped off (the box can't score, so a credit-out now fails cleanly instead of hanging).

- **Live JobAdder feed** — the operator's JobAdder candidate data (~13k profiles) is available live and on demand, drawn from the operator's own JobAdder account, and kept current in the platform.

### Known constraints / debt

- **Two parallel search systems** — now largely converged onto one durable, resumable multi-source run (job-scoped, live-streaming). The legacy job-search path is mostly retired; final dead-code cleanup is the remaining item — see roadmap §3.
- SEEK candidates are **snippet-level** (name/role/company/location); full-profile enrichment is a separate, credit-gated action (scraper exists, not auto-fired). *(Fixed 2026-06-15: "Score all" now scores these thin LinkedIn/SEEK finds as a low-confidence estimate instead of silently skipping them.)*
- ~15.4k candidates, **731 scored** — large unscored backlog. Scoring is **manual + on-demand** by design (recruiter clicks *Score all* or per-card) and cost-capped; finds are not auto-scored on import/discovery.
- **No AI fallback** — production is 100% Claude-dependent (no working local model on Railway or the box). If the Anthropic credit balance runs out or Claude is unavailable, all AI — scoring **and** insight extraction — hard-fails with no graceful degradation. Operationally: keep the Anthropic balance topped up.
- Single scraper worker (no redundancy); the box must stay online for live discovery.
- Scraper still authenticates via a shared secret (org-binding token cutover incomplete).

---

## 3. Future implementations (roadmap)

---

*The authoritative, sequenced plan to take the app to feature-complete + self-running lives in `docs/completion-plan.md` (7 phases, built from a verified 6-dimension audit). The items below are the strategic summary.*

**A. Search architecture consolidation** *(planned, documented in `docs/search-architecture-plan.md`)* — collapse the two parallel search systems into one **federated dispatcher + pluggable source adapters** (pool/LinkedIn/SEEK/+future), one durable result model, and a **live-streaming results UI** (results appear as they're scraped). Makes adding a new source trivial and removes the current fragility. ~6 phases, each independently shippable.

**B. SEEK full-profile enrichment** — one-click, credit-aware deep-fetch of a chosen SEEK candidate's full profile (the scraper exists; needs wiring + a credit guard).

**C. Scoring at scale** — score the unscored backlog; background/batch scoring; richer re-score-on-change.

**D. Talent flywheel — later phases** — the **read-path (insight-ranked search) is now live + backfilled (812 profiles)**. Remaining: background discovery that continuously enriches the pool (cost-gated); profile freshness / re-fetch of stale profiles; raising the insight coverage as the pool grows.

**E. Candidate UX polish** — **shipped** (email/screening/file-count surfacing, import-batch grouping, source badges). Remaining: candidate-card density declutter.

**F. Hardening** — complete the scraper org-binding token cutover; multi-worker / redundancy if scrape volume grows; observability.

Priority order is the operator's call; A unblocks the most future work.

## 4. Running costs per month

| Item                                       | Cost / month                            | Basis                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|--------------------------------------------|-----------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Claude API (Anthropic)</b>              | <b>~\$15–40</b>                         | Mostly cheap Haiku (parsing, snippet scoring, insight extraction). With Sonnet scoring full profiles (~\$0.05/score) cost rises with scoring volume; <b>hard-capped at \$5/day</b> (≈ \$150/mo ceiling) by the spend guard. Realistic steady-state ~\$20–40. One-off: the talent-flywheel insight backfill (812 profiles, Haiku) cost <b>~\$4</b> total. ⚠️ No fallback — if this balance hits zero, all AI stops (see constraints §2). |
| <b>Railway</b> (app + Postgres + egress)   | <b>\$25</b>                             | Hosting: always-on Next app + Postgres (5 GB volume) + bandwidth. Operator-confirmed figure.                                                                                                                                                                                                                                                                                                                                            |
| <b>CV storage</b> (S3 / t3.storageapi.dev) | <b>~\$1–3</b>                           | ~13.5k encrypted CVs (~8–12 GB est.) at S3-compatible rates + light egress.                                                                                                                                                                                                                                                                                                                                                             |
| <b>Mini-PC scraper</b> (electricity)       | <b>~\$3–6</b>                           | Low-power i3 mini-PC 24/7 at NZ power rates. Internet = existing home line; Tailscale = free tier.                                                                                                                                                                                                                                                                                                                                      |
| <b>Domain</b>                              | <b>\$0</b>                              | Railway subdomain ( *.up.railway.app ). Custom domain would add ~\$1–2/mo.                                                                                                                                                                                                                                                                                                                                                              |
| <b>Platform infra subtotal</b>             | <b>~\$45–75 / month</b>                 | Everything RecruitMe itself costs to run (recurring).                                                                                                                                                                                                                                                                                                                                                                                   |
| <b>Mini-PC hardware</b> (one-off)          | <b>\$150 once</b>                       | The mini-PC itself — already purchased; one-time capital cost, not recurring (≈ \$4–6/mo amortised over 2–3 yrs).                                                                                                                                                                                                                                                                                                                       |
| <b>SEEK Talent Search</b>                  | <i>operator's existing subscription</i> | External recruiting tool (the box logs into the operator's SEEK account). Significant but pre-existing; per-search/profile credits apply — the app is built to be credit-safe.                                                                                                                                                                                                                                                          |
| <b>LinkedIn</b>                            | <i>operator's existing account</i>      | Scraper uses the operator's logged-in LinkedIn; no extra app cost.                                                                                                                                                                                                                                                                                                                                                                      |

**Bottom line:** the RecruitMe *platform* runs at roughly **\$45–75/month** in direct infrastructure + AI cost (plus a one-off **\$150** mini-PC, already purchased), on top of the operator's existing SEEK/LinkedIn/JobAdder subscriptions. The single largest and most variable line is the Claude API, which is bounded by the daily spend cap and scales with how much scoring is done.

## 5. Tech-stack summary (for a developer handoff)

---

- **Frontend/Backend:** Next.js 15 (App Router), React 19, TypeScript, Tailwind.
- **Data:** PostgreSQL 18 + Prisma; Postgres full-text search (FTS) for the talent pool.
- **AI:** Anthropic SDK (Claude Haiku 4.5 + Sonnet 4.6); structured scoring pipeline with caching + cost tracking + spend caps.
- **Scraping:** Node/TypeScript worker on a mini-PC, Patchright (patched Chromium), one persistent session per platform; polls the app's job queue (Postgres-backed `ScrapeJob` , priority + dedup).
- **Infra:** Railway (CI real-DB gate on deploy), Tailscale (box networking), S3-compatible blob storage (encrypted CVs).
- **Repo:** GitHub `baeley-canning/RecruitMe` , deploy = push to `main` .